

▼ PHASE 3 PROJECT: SYRIATEL CUSTOMER CHURN

- Business Understanding
- Data Loading and Initial Inspection
- Data Cleaning
- Exploratory Data Analysis
- Prepare for Data Modelling
- Build and Evaluate Models
- Model Comparison
- Feature Importance(for best model if tree based)
- Business Impact
- Recommendations

▼ 1. BUSINESS UNDERSTANDING

"" Business Problem: SyriaTel is a telecommunications company facing customer churn issues. Churn refers to customers leaving the service, which directly impacts revenue.

Stakeholder: SyriaTel Retention Department Business Goal: Identify customers likely to churn 30-60 days in advance Success Metric: Reduce churn rate by 15% within 6 months

.....

▼ 2. DATA LOADING AND INITIAL INSPECTION

```
!pip install kagglehub
```

```
Requirement already satisfied: kagglehub in c:\users\user\anaconda3\lib\site-pac
Requirement already satisfied: packaging in c:\users\user\anaconda3\lib\site-pac
Requirement already satisfied: pyyaml in c:\users\user\anaconda3\lib\site-packag
Requirement already satisfied: requests in c:\users\user\anaconda3\lib\site-pack
Requirement already satisfied: tqdm in c:\users\user\anaconda3\lib\site-packages
Requirement already satisfied: charset_normalizer<4,>=2 in c:\users\user\anacond
Requirement already satisfied: idna<4,>=2.5 in c:\users\user\anaconda3\lib\site-
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\user\anaconda3\lib
```

```
Requirement already satisfied: certifi>=2017.4.17 in c:\users\user\anaconda3\lib  
Requirement already satisfied: colorama in c:\users\user\anaconda3\lib\site-pack
```

```
import kagglehub  
  
# Download latest version  
path = kagglehub.dataset_download("becksddf/churn-in-telecoms-dataset")  
  
print("Path to dataset files:", path)
```

```
Warning: Looks like you're using an outdated `kagglehub` version (installed: 0.3  
Path to dataset files: C:\Users\USER\.cache\kagglehub\datasets\becksddf\churn-in
```

```
# Import necessary libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.model_selection import train_test_split, GridSearchCV  
from sklearn.preprocessing import StandardScaler  
from sklearn.linear_model import LogisticRegression  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier  
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score  
from imblearn.over_sampling import SMOTE  
import warnings
```

```
#Load the data  
df = pd.read_csv('syriatel_churn.csv')  
print(f"Dataset shape: {df.shape}")
```

```
Dataset shape: (3333, 21)
```

```
df.head()
```

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls
0	KS	128	415	382-4657		no	yes	25	265.1
1	OH	107	415	371-7191		no	yes	26	161.6
2	NJ	137	415	358-1921		no	no	0	243.4
3	OH	84	408	375-9999		yes	no	0	299.4
4	OK	75	415	330-6626		yes	no	0	166.7

5 rows × 21 columns

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   state            3333 non-null    object 
 1   account length   3333 non-null    int64  
 2   area code         3333 non-null    int64  
 3   phone number     3333 non-null    object 
 4   international plan 3333 non-null    object 
 5   voice mail plan  3333 non-null    object 
 6   number vmail messages 3333 non-null    int64  
 7   total day minutes 3333 non-null    float64
 8   total day calls   3333 non-null    int64  
 9   total day charge  3333 non-null    float64
 10  total eve minutes 3333 non-null    float64
 11  total eve calls   3333 non-null    int64  
 12  total eve charge  3333 non-null    float64
 13  total night minutes 3333 non-null    float64
 14  total night calls  3333 non-null    int64  
 15  total night charge 3333 non-null    float64
 16  total intl minutes 3333 non-null    float64
 17  total intl calls   3333 non-null    int64  
 18  total intl charge  3333 non-null    float64
 19  customer service calls 3333 non-null    int64  
 20  churn             3333 non-null    bool  
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

df.shape

(3333, 21)

df.describe()

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.56230
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.25943
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.00000
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.43000
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.50000
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.79000
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.64000

3. DATA CLEANING

```
# Make a copy of the original dataframe for preprocessing
df_clean = df.copy()
```

```
# Remove phone number (unique identifier)
if 'phone number' in df_clean.columns:
    df_clean = df_clean.drop('phone number', axis=1)
    print("✓ Dropped 'phone number' column")
```

✓ Dropped 'phone number' column

```
# Handle missing values
missing_before = df_clean.isnull().sum().sum()
if missing_before > 0:
    print(f"\nHandling {missing_before} missing values...")
    # Fill numeric with median
    for col in df_clean.select_dtypes(include=[np.number]).columns:
        if df_clean[col].isnull().sum() > 0:
            df_clean[col] = df_clean[col].fillna(df_clean[col].median())

    # Fill categorical with mode
    for col in df_clean.select_dtypes(include=['object']).columns:
        if df_clean[col].isnull().sum() > 0:
```

```
df_clean[col] = df_clean[col].fillna(df_clean[col].mode()[0])

print("✓ Missing values filled")
```

```
# Check for duplicates
duplicates = df_clean.duplicated().sum()
if duplicates > 0:
    print(f"\nRemoving {duplicates} duplicate rows...")
    df_clean = df_clean.drop_duplicates()
    print("✓ Duplicates removed")

print(f"\nFinal dataset shape: {df_clean.shape}")
```

```
Final dataset shape: (3333, 20)
```

▼ 4. EXPLORATORY DATA ANALYSIS


```
# First, get the churn counts from cleaned data
churn_counts = df_clean['churn'].value_counts()
print(f"Churn distribution in cleaned data: {dict(churn_counts)}")
```

```
Churn distribution in cleaned data: {False: 2850, True: 483}
```

```
# Visual 1: Churn distribution
plt.figure(figsize=(12, 4))


plt.subplot(1, 2, 1)
plt.pie(churn_counts, labels=['Not Churned', 'Churned'],
        autopct='%.1f%%', colors=['lightblue', 'salmon'])
plt.title('Customer Churn Distribution')

plt.subplot(1, 2, 2)
plt.bar(['Not Churned', 'Churned'], churn_counts, color=['lightblue', 'salmon'])
plt.title('Customer Churn Count')
plt.ylabel('Number of Customers')
plt.show()
```




```
# Visual 2: Churn by international plan
if 'international plan' in df_clean.columns:
    plt.figure(figsize=(8, 4))
    churn_by_plan = pd.crosstab(df_clean['international plan'], df_clean['churn'])
    churn_by_plan.plot(kind='bar')
    plt.title('Churn Rate by International Plan')
    plt.ylabel('Percentage')
    plt.legend(['Not Churned', 'Churned'])
    plt.show()
```


<Figure size 800x400 with 0 Axes>



```
# Visual 3: Churn by customer service calls
if 'customer service calls' in df_clean.columns:
    plt.figure(figsize=(10, 4))
    service_calls_churn = df_clean.groupby('customer service calls')['churn'].n
    service_calls_churn.plot(kind='bar', color='coral')
    plt.title('Churn Rate by Customer Service Calls')
    plt.xlabel('Customer Service Calls')
    plt.ylabel('Churn Rate (%)')
    plt.axhline(y=df_clean['churn'].mean() * 100, color='red', linestyle='--',
                label=f'Overall Rate ({df_clean["churn"].mean()*100:.1f}%)')
    plt.legend()
    plt.show()
```



```
# Visual 4: Correlation heatmap
plt.figure(figsize=(10, 8))
numeric_df = df_clean.select_dtypes(include=[np.number])
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', center=0, fmt='.2f')
plt.title('Feature Correlation Heatmap')
plt.show()
```



5. PREPARE DATA FOR MODELING

```
# Create copy for modeling
df_model = df_clean.copy()

# Encode categorical variables
cat_cols = df_model.select_dtypes(include=['object']).columns
for col in cat_cols:
    if df_model[col].nunique() == 2:
        # Binary encoding
        df_model[col] = df_model[col].map({'yes':1, 'no':0, 'Yes':1, 'No':0})
        print(f"✓ Binary encoded '{col}'")
    else:
        # One-hot encoding
        dummies = pd.get_dummies(df_model[col], prefix=col, drop_first=True)
        df_model = pd.concat([df_model.drop(col, axis=1), dummies], axis=1)
        print(f"✓ One-hot encoded '{col}'")
```

✓ One-hot encoded 'state'
✓ Binary encoded 'international plan'
✓ Binary encoded 'voice mail plan'

```
# Features and target
X = df_model.drop('churn', axis=1)
y = df_model['churn']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# SMOTE to balance classes
smote = SMOTE(random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)
print(f"\nBefore SMOTE: {np.bincount(y_train)}")
print(f"After SMOTE: {np.bincount(y_train_bal)})")
```

Before SMOTE: [2280 386]
After SMOTE: [2280 2280]

```
# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)
print(f"\nScaled training data: {X_train_scaled.shape}")
print("✓ Data preparation complete")
```

Scaled training data: (4560, 68)
✓ Data preparation complete

6. BUILD AND EVALUATE MODELS

```
# evaluation function (no external accuracy_score needed)
def evaluate_model(model, X_train, X_test, y_train, y_test, model_name):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # Compute accuracy directly
    acc = (y_pred == y_test).mean()
    recall = recall_score(y_test, y_pred)

    # AUC if available
    auc = None
    if hasattr(model, 'predict_proba'):
        y_proba = model.predict_proba(X_test)[:, 1]
        auc = roc_auc_score(y_test, y_proba)
    print(f"\n{model_name}")
    print("-" * 30)
    print(f"Test Accuracy: {acc:.4f}")
    print(f"Test Recall: {recall:.4f}")
    if auc:
        print(f"Test AUC: {auc:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred, target_names=['Not Churned', 'Churned']))

return {'model': model, 'accuracy': acc, 'recall': recall, 'auc': auc, 'y_pred': y_pred}
```

```
# ----- Model 1: Logistic Regression (baseline) -----
lr = LogisticRegression(random_state=42, max_iter=1000)
lr_res = evaluate_model(lr, X_train_scaled, X_test_scaled, y_train_bal, y_test, model_name='Logistic Regression (baseline)')

# ----- Model 2: Logistic Regression (tuned) -----
lr_tuned = LogisticRegression(random_state=42, max_iter=1000, class_weight='balanced')
lr_tuned_res = evaluate_model(lr_tuned, X_train_scaled, X_test_scaled, y_train_bal, y_test, model_name='Logistic Regression (tuned)')

# ----- Model 3: Random Forest -----
rf = RandomForestClassifier(random_state=42, n_estimators=100, class_weight='balanced')
rf_res = evaluate_model(rf, X_train_scaled, X_test_scaled, y_train_bal, y_test, model_name='Random Forest')

# ----- Model 4: Decision Tree (simple) -----
dt = DecisionTreeClassifier(random_state=42, max_depth=5)
dt_res = evaluate_model(dt, X_train_scaled, X_test_scaled, y_train_bal, y_test, model_name='Decision Tree (simple)')

# Collect results
results = {
    'Logistic Regression (baseline)': lr_res,
    'Logistic Regression (tuned)': lr_tuned_res,
    'Random Forest': rf_res,
```

```
'Decision Tree': dt_res  
}
```

1. Logistic Regression (baseline)

```
-----  
Test Accuracy: 0.8576  
Test Recall: 0.2887  
Test AUC: 0.7898
```

Classification Report:

	precision	recall	f1-score	support
Not Churned	0.89	0.95	0.92	570
Churned	0.52	0.29	0.37	97
accuracy			0.86	667
macro avg	0.70	0.62	0.65	667
weighted avg	0.83	0.86	0.84	667

2. Logistic Regression (tuned)

```
-----  
Test Accuracy: 0.8576  
Test Recall: 0.3196  
Test AUC: 0.7932
```

Classification Report:

	precision	recall	f1-score	support
Not Churned	0.89	0.95	0.92	570
Churned	0.52	0.32	0.39	97
accuracy			0.86	667
macro avg	0.70	0.63	0.66	667
weighted avg	0.84	0.86	0.84	667

3. Random Forest

```
-----  
Test Accuracy: 0.8936  
Test Recall: 0.5876  
Test AUC: 0.8645
```

Classification Report:

	precision	recall	f1-score	support
Not Churned	0.93	0.95	0.94	570
Churned	0.65	0.59	0.62	97
accuracy			0.89	667
macro avg	0.79	0.77	0.78	667
weighted avg	0.89	0.89	0.89	667

4. Decision Tree

```
-----  
Test Accuracy: 0.8621  
Test Recall: 0.6392  
Test AUC: 0.7791
```

8. MODEL COMPARISON


```
comp_df = pd.DataFrame([  
    {'Model': name, 'Accuracy': res['accuracy'], 'Recall': res['recall'], 'AUC':  
        for name, res in results.items()  
    ]).sort_values('Recall', ascending=False)  
  
print(comp_df.to_string(index=False))
```

Model	Accuracy	Recall	AUC
Decision Tree	0.862069	0.639175	0.779146
Random Forest	0.893553	0.587629	0.864478
Logistic Regression (tuned)	0.857571	0.319588	0.793163
Logistic Regression (baseline)	0.857571	0.288660	0.789799

```
# Select best model based on recall  
best_idx = comp_df['Recall'].idxmax()  
best_model_name = comp_df.loc[best_idx, 'Model']  
best_model = results[best_model_name]['model']  
best_y_pred = results[best_model_name]['y_pred']  
print(f"\n\n{checkmark} Best model: {best_model_name} (Recall = {comp_df.loc[best_idx, 'R}
```

{checkmark} Best model: Decision Tree (Recall = 0.6392)

```
# Show confusion matrix only for the best model  
cm = confusion_matrix(y_test, best_y_pred)  
plt.figure(figsize=(5,4))  
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',  
            xticklabels=['Pred Not Churned', 'Pred Churned'],  
            yticklabels=['Actual Not Churned', 'Actual Churned'])  
plt.title(f'Confusion Matrix - {best_model_name}')  
plt.ylabel('True')  
plt.xlabel('Predicted')  
plt.show()
```



9. FEATURE IMPORTANCE (for best model if tree-based)

```
if hasattr(best_model, 'feature_importances_'):
    imp = best_model.feature_importances_
    feat_names = X.columns
    imp_df = pd.DataFrame({'feature': feat_names, 'importance': imp}).sort_values('importance', ascending=False)
    print("\nTop 10 important features:")
    print(imp_df.head(10).to_string(index=False))

    # Plot
    plt.figure(figsize=(10,5))
    top10 = imp_df.head(10)
    plt.barh(range(len(top10)), top10['importance'])
    plt.yticks(range(len(top10)), top10['feature'])
    plt.xlabel('Importance')
    plt.title(f'Top 10 Feature Importances ({best_model_name})')
    plt.gca().invert_yaxis()
    plt.tight_layout()
    plt.show()
else:
    print("Feature importance not available for this model.")
```

Top 10 important features:

	feature	importance
1	total day charge	0.284936
2	customer service calls	0.261715
3	voice mail plan	0.103382
4	total eve charge	0.075338
5	total intl calls	0.067741
6	total day minutes	0.059565
7	international plan	0.058011
8	total intl charge	0.052572
9	total night charge	0.010891
10	state_WY	0.006641

Top 10 Feature Importances (Decision Tree)

